

Exercise

Using an autoencoder to pre-train an image classifier: split the MNIST data to training and test sets. 1) Train a deep denoising autoencoder on the full training set. Check that the autoencoder works well (i.e. check if the images are fairly well reconstructed), and visualize the low-level features. Visualize images that most activate each neuron in the coding layer. 2) Build a classification deep neural network, by reusing the lower layers of the autoencoder. Train it using only 10% of the training set. How does the performance of this classifier compare to the same classifier trained on the full training set?

Approach:

- Developed Encoder-Decoder model
- Trained the model with MNIST data
 - For this, we don't need labelled data as we only want to project the higher dimensional image to non-linear lower-dimensional data
- And the error matrix we can use can be MSE
- Lower level filters are visualized
 - conv2D_1
 - conv2D_2
- CNN based supervised model is trained with using the Autoencoder model's early layers
 - Using only 10% of the training dataset
- The same model was trained with the full dataset
- The accuracy and other evaluation matrix are compared

Results:

- Structure of Encoder-Decoder model:

Model: "encoder"

Layer (type)	Output Shape	Param #
encoder_input (InputLayer)	(None, 28, 28, 1)	0
conv2d_1 (Conv2D)	(None, 14, 14, 32)	320
conv2d_2 (Conv2D)	(None, 7, 7, 64)	18496
flatten_1 (Flatten)	(None, 3136)	0
latent_vector (Dense)	(None, 32)	100384

Total params: 119,200
Trainable params: 119,200
Non-trainable params: 0


Model: "decoder"

Layer (type)	Output Shape	Param #
decoder_input (InputLayer)	(None, 32)	0
dense_1 (Dense)	(None, 3136)	103488
reshape_1 (Reshape)	(None, 7, 7, 64)	0
conv2d_transpose_1 (Conv2DTr	(None, 14, 14, 64)	36928
conv2d_transpose_2 (Conv2DTr	(None, 28, 28, 32)	18464
conv2d_transpose_3 (Conv2DTr	(None, 28, 28, 1)	289
decoder_output (Activation)	(None, 28, 28, 1)	0

Total params: 159,169
Trainable params: 159,169
Non-trainable params: 0

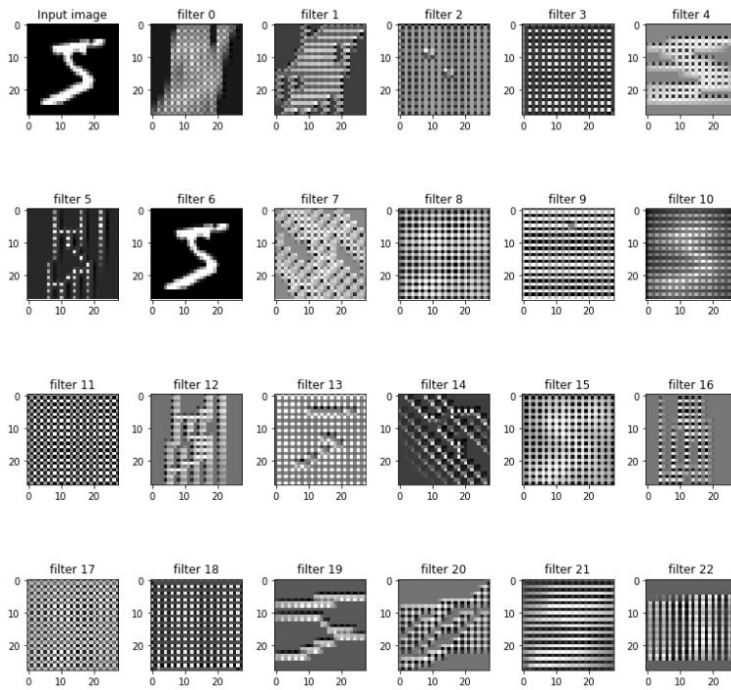
Denoising on "noisy" test images:

Original images in top rows, Corrupted Input in middle rows, Denoised Input in third rows




- Visualize images that most activate each neuron in the encoding layer

Conv2D_1 output:



Conv2D_2 output:



Part 2 :

Using the Encoder model as a lower level feature detector and train the model:

Model: "sequential_2"

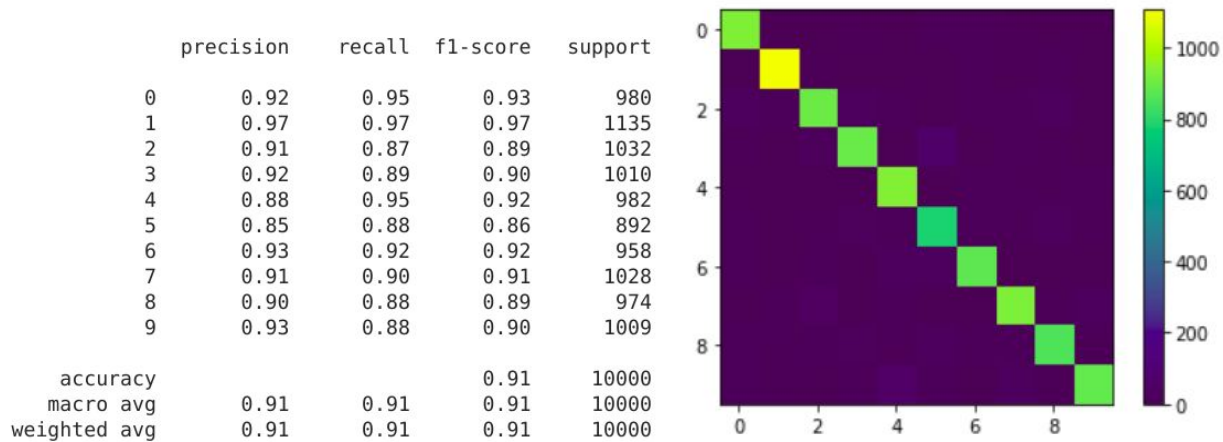
Layer (type)	Output Shape	Param #
encoder (Model)	(None, 32)	119200
dense_5 (Dense)	(None, 16)	528
dense_6 (Dense)	(None, 16)	272
dense_7 (Dense)	(None, 10)	170

Total params: 120,170

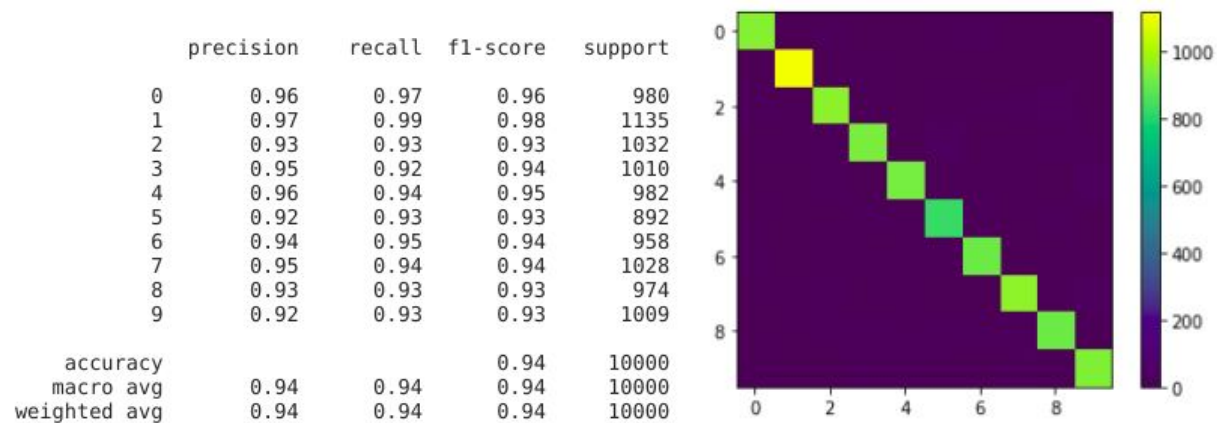
Trainable params: 970

Non-trainable params: 119,200

Result with 10% of data:



Result with 100% of data:



Conclusion:

- We can say that just by using the initial layers of AE, the model is able to fit the dataset with really very fewer data without losing accuracy.
- AE is able to denoise the images very effectively
 - The images contain salt and pepper noise the extent to which humans are not able to recognize the digit, yet AE performs well.